EchoScribe: Evidence-Linked, Independently Verified Notes for Therapy and Mentoring

Aravind Krishnan S, Parvathy Krishna M, Angisha B, Lloyd Sebastian
Department of Artificial Intelligence and Data Science
Jyothi Engineering College, Cheruthuruthy, Kerala, India
{aravindkrishnans.ad23, parvathykrishnam.ad23, angishab.ad23, lloydsebastian.ad23}@jecc.ac.in

Jayakrishnan A
Department of Artificial Intelligence and Data Science
Jyothi Engineering College, Cheruthuruthy, Kerala, India
jayakrishnana@jecc.ac.in

***Abstract*—Therapists and academic mentors spend much of their working time writing session records. Language models can draft these records from a recording, but a draft may state things the session never contained, and a single cloud model is a single point of failure for quota, outage and data-governance reasons. We present EchoScribe, an open system that turns a recorded two-person session into a SOAP (therapy) or GROW (mentoring) note in which every sentence cites the transcript turns that support it. A model other than the writer, from another vendor when one is configured, checks each sentence against its cited turns, and notes that fail are regenerated on another model. An embedded multi-provider router keeps every processing role available through quota exhaustion and outages, and speaker separation runs locally on the original recording, with each word assigned to a speaker by its timestamp. On the 40 test dialogues of ACI-Bench, citing evidence halved unsupported content against a single-pass note from the same writer (11.2% to 5.0% of sentences) without lowering overlap with reference notes, and the verifier caught 98% of planted errors. On 57 mock consultations, speaker-attributed word error was 23.8% with word timestamps and 53.0% without. Clinician or mentor identification needed no enrolment and was correct in 57 of 57 consultations and 40 of 40 counselling dialogues. In fault-injection simulations, a last-known-good routing strategy cut 95th-percentile latency from 31.7 s to 9.1 s. We state plainly what remains unvalidated.**

***Index Terms*—clinical documentation, hallucination detection, speaker diarization, large language models, LLM routing, SOAP notes, GROW coaching**

I. Introduction

Writing session notes is one of the most time-consuming non-clinical tasks of mental health work and of academic mentoring. A SOAP note records the client's account, the clinician's observations, an assessment and a plan [1]; a GROW note records a mentee's goal, current reality, options and way forward [2]. Documentation extends the working day well beyond the session [3], and ambient scribes that draft notes from recorded encounters are now in clinical use and under randomized evaluation [4], [5].

Two gaps limit what such systems can be trusted with. A note may contain statements the conversation never supported [6], [7], and the reviewer cannot tell which sentence rests on what. And published scribes depend on one model from one vendor, so a quota, an outage or a retired model stops documentation altogether; neither gap is addressed for counselling or mentoring [8]–[10]. EchoScribe's contributions are:

1) **Evidence-linked notes with independent verification.** The writer cites turn numbers for every sentence; a model other than the writer, preferably from another vendor, labels each sentence supported, partially supported or unsupported; notes that fail a threshold are regenerated on another model, and flagged sentences can be repaired from the transcript (Section III-C).
2) **An embedded multi-provider router** serving every processing role from a combo of targets with a routing strategy, per-key exponential cooldown and a provider-wide breaker, adapted from the open-source OmniRoute gateway [11] (Section III-A).
3) **Local speaker attribution without enrolment:** pyannote diarization of the original recording, words assigned to turns by their timestamps, and role resolution from eight conversation statistics and a model check (Section III-B).
4) **A reproducible evaluation** on public data (ACI-Bench, AnnoMI, PriMock57, LibriSpeech-based simulations) of the router under faults, note grounding against a single-pass baseline, verifier accuracy on planted errors, role identification, and speech under noise (Section IV).

II. Related Work

Ambient scribes and note generation. Note generation from doctor–patient dialogue has been benchmarked on ACI-Bench [10] and in MEDIQA-Chat [12]; PriMock57 adds audio of mock consultations [13] and AnnoMI counselling dialogues [14]. Adapted models can match experts on clinical summarisation [15], [16], evaluation frameworks combine metrics, reviewers and simulation [8], and deployed systems report physician review [9]. For mental health, Adhikary et al. benchmarked summarisation of counselling sessions [17]. None of these produces mentoring (GROW) notes, links each sentence to its evidence, or survives its vendor failing.

Hallucination and verification. Hallucination is a recurring failure of generated text [6], and clinical safety frameworks annotate hallucinations and omissions [7]. Remedies include citation-constrained generation [18], atomic-fact scoring [19], self-consistency sampling [20] and self-verification [21]. EchoScribe combines citation-constrained writing with sentence-level verification by an *independent* model, so that the checker does not share the writer's failure modes, and uses the verdicts for review and to trigger regeneration.

Speech. Diarization combines activity detection, speaker embeddings and clustering [22]: ECAPA-TDNN [23] with spectral clustering [24] is a strong baseline, while pyannote's powerset model [25] and Sortformer [26] learn segmentation directly; WhisperX aligns recognition with diarization [27], and speaker roles in clinical conversations have been studied [28]. SepFormer [29] in SpeechBrain [30], trained on WHAM! [31], is one of many enhancement models [32]; enhancement artefacts can hurt recognition [33], which we test rather than assume.

Routing. RouteLLM routes between a strong and a weak model [34], FrugalGPT cascades models to cut cost [35], and gateways such as OmniRoute [11] provide combo routing with cooldowns. EchoScribe embeds this per processing role and adds verification as a routing signal: a note is escalated when its content fails, not only when a call does.

III. System Design

Fig. 1 shows the pipeline. A browser client records or uploads audio; an Express.js server coordinates every model call; Supabase stores records under per-user row-level security, with therapy and mentoring records separated by a mode column that every query must state. Five roles call language or speech models: STT, DIARIZE, ROLE_LABEL, NOTE_SYNTH and VERIFY.

A. Embedded multi-provider router

Each role is served by a *combo*: an ordered list of targets $\langle$vendor, model, keys$\rangle$ and a routing strategy. Vendors are thin adapters; Gemini [36], Groq-hosted gpt-oss-120b [37] and Whisper [38], Deepgram and a local pyannote service ship as adapters, and OpenAI-compatible APIs are added by configuration. Several keys for one model form a pool whose keys take turns. The strategy orders targets per request: *priority* (declared order; the next target is used only when every key of the current one is cooling down or retired, or the model itself fails), *round-robin*, *random*, *strict-random*, *least-used*, and *lkgp* (the last target that succeeded goes first).

Failures are classified. A key that fails authentication is retired; a rate-limited or failing key enters an exponential cooldown of $5 \cdot 2^{k-1}$ s (capped at 5 min) after its k-th consecutive failure, reset on success; a vendor with 15 failures within 30 min is skipped for 10 min (2 in 5 min for the local service); an HTTP 400 is treated as specific to the request and moves on without cooling the key. These mechanisms and constants are adapted from OmniRoute [11]. When every target fails the API returns HTTP 503 with the attempts. Keys stay in the user's browser unless the user syncs one, which is then stored encrypted.

B. Speech processing and role resolution

The recording is not denoised: SepFormer enhancement, our first design, raised word error rates (Section IV), as reported for other front ends [33]. DIARIZE runs pyannote community-1 [25], [39] locally on the original 16 kHz recording, on a GPU when present, with two speakers and non-overlapping turns; it falls back to cloud audio diarization, then to splitting the transcript with a language model. STT transcribes the original recording; when the recogniser returns word timestamps (Whisper, Deepgram), each word goes to the turn that contains it [27], otherwise text is spread over turns by duration. A denoised copy (FFmpeg) is made only for listening again.

Roles are resolved without voice enrolment. For each speaker the system computes eight statistics: speaking time, word count and mean turn length, questions, professional vocabulary, guiding phrases, personal and emotional phrases, number of turns, and who spoke first. ROLE_LABEL receives the transcript with these statistics and must assign exactly one professional and one client; invalid answers move to the next target, and a deterministic default is used only if every target fails.

C. Evidence-linked notes, verification and repair

The labelled transcript is numbered into turns $T_1 \dots T_n$. NOTE_SYNTH must end every sentence of the note sections, and of the risk notes, with the turns that support it ("…poor sleep for two weeks [T4, T6]"), write one fact per sentence, and write "Not discussed" rather than an uncited sentence. Citations are parsed into per-sentence evidence and removed from the displayed text. VERIFY receives each sentence with only its cited turns and labels it *supported*, *partial* (adds a detail the turns lack) or *unsupported*. Verifiers are ordered another vendor first, then another model of the writer's vendor, and never the writer's own model. If the unsupported fraction exceeds $\tau = 0.2$, the note is regenerated with the writer's model moved to the end of the combo, re-verified, and the draft with fewer unsupported sentences is kept.

Two further steps work at sentence level. *Repair*: each flagged sentence goes back to a writer model with the transcript and the verifier's reason to be rewritten with citations, or removed if the transcript does not support it; rewrites are verified again. *Two-judge agreement*: a flag stands only if a second judge, chosen to differ from the first and from the writer, also flags it. They are implemented in the released code but not evaluated in this paper.The interface shows each sentence's evidence on hover and marks partial and unsupported sentences (Fig. 2).

D. Risk screen with a second opinion

In therapy mode the writer also screens for suicidal ideation and self-harm; in mentoring mode, for burnout and severe distress. A second model, chosen like the verifier, repeats the

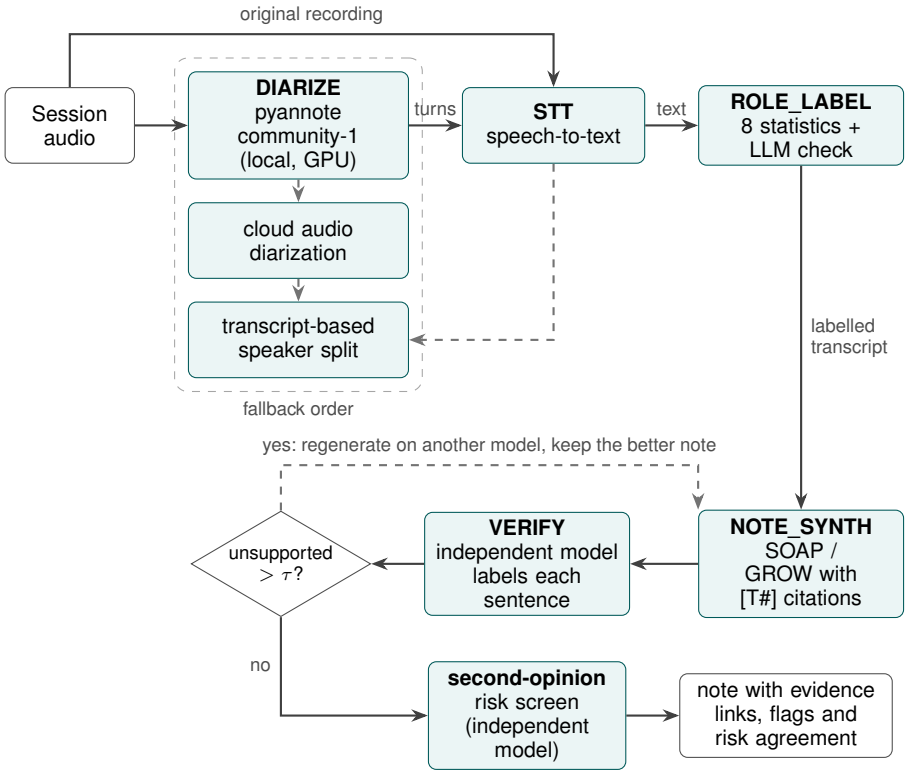


Fig. 1. EchoScribe pipeline. Shaded boxes are processing roles, each served by a combo of the embedded router (Section III-A). Speaker separation falls back from local pyannote to cloud audio diarization to a transcript-based split (dashed). The note writer cites transcript turns; an independent model (another vendor, else another model) verifies each sentence, and notes over the unsupported threshold τ are regenerated on another model. A second model screens independently for risk.

screen independently and cites its evidence turns. The flag is raised if either model raises it, favouring sensitivity, and disagreement is shown as "models disagree — review the transcript". The screen is a model judgement, not a validated instrument such as the C-SSRS or PHQ-9 [40], [41], and the interface states that the clinician remains responsible for risk assessment.

IV. Evaluation

Every number below is generated from result files by scripts in the public repository (`eval/`); only free-tier model access was used.

A. Setup

Data. Notes are evaluated on the 40 test dialogues of ACI-Bench (test1, CC BY 4.0) [10], whose doctor–patient visits have reference notes, and, without reference notes, on 40 high-quality motivational-interviewing dialogues of AnnoMI [14] drawn with a fixed seed over 32 topics. Real conversational speech comes from PriMock57 [13]: 57 mock primary-care consultations (8.64 h) acted by clinicians and employees, with manual transcripts (CC BY 4.0). Controlled noise tests use 12 simulated two-person conversations built from LibriSpeech test-clean [42] (ten alternating utterances, 0.3–0.9 s gaps, exact reference turns), clean and with babble plus pink noise at 5 dB and 0 dB SNR. **Models and metrics.** Gemini models [36] are used in rotation across three keys; gpt-oss-120b [37] and Whisper large-v3 [38] are served by Groq. We report ROUGE [43] and BERTScore [44] against reference notes, WER after text normalisation, DER with a 0.25 s collar [45], and cpWER, the speaker-attributed WER of CHiME-6 [46]. Intervals are Wilson 95%; differences are tested with a paired bootstrap (10,000 resamples).

B. Router under injected faults

We drive the router's own code with simulated vendors on a virtual clock: one hour of Poisson traffic at 10 requests/min against a preferred vendor (two keys), a fallback vendor (two keys) and a second fallback, in three scenarios: a per-key quota below the offered load, a 15-minute outage of the preferred vendor, and a *mixed* case with a revoked key, a quota and a fallback that times out 20% of the time. We compare walking the chain with no memory, the fixed 60 s suspension of our earlier design, and three strategies (5 seeds, Table I).

Where a healthy fallback exists no policy changes success much, so differences lie in wasted calls and latency. Last-known-good stops returning to a vendor that keeps timing out and cuts mixed-scenario p95 latency from 31.7 s to 9.1 s (difference $[-28.3, -11.7]$ s, $p < 0.001$, faster in all 5 seeds). Under a per-key quota, round-robin and LKGP wasted 0.014 and 0.033 failed calls per request against 0.227 for priority; exponential cooldown alone is not uniformly better, because

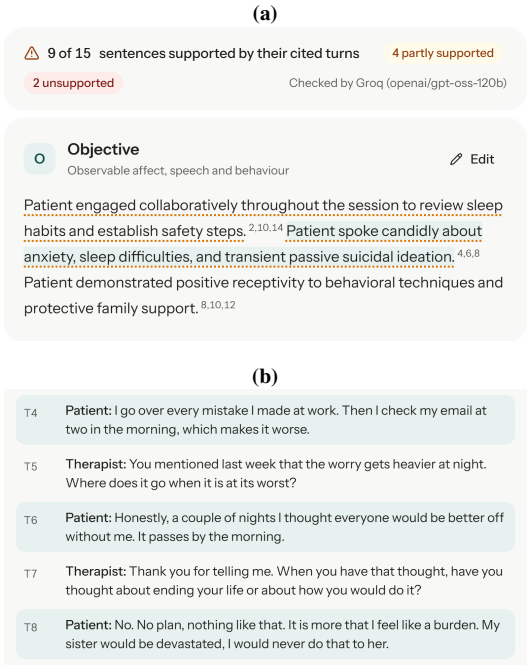


Fig. 2. EchoScribe's session summary for a scripted example session (invented transcript; real pipeline output). (a) The verification bar counts sentences the independent verifier (here gpt-oss-120b, checking a Gemini-written note) judged supported, partly supported or unsupported; in the Objective section, dotted underlines mark partly supported sentences and superscripts give cited turns. (b) Hovering the second Objective sentence highlights the turns it cites (T4, T6, T8) in the transcript.

TABLE I
ROUTER IN THE MIXED-FAULT SCENARIO (MEAN OF 5 SEEDS): SUCCESS RATE, FAILED PROVIDER CALLS PER REQUEST, 95TH-PERCENTILE LATENCY.

Policy	Success	Failed calls/req.	p95 (s)
No memory	99.93%	0.497	31.8
Fixed 60 s	98.14%	0.147	31.5
Priority	99.86%	0.258	31.7
Round-robin	100.00%	0.092	31.4
LKGP	100.00%	0.097	9.1

its 5 s first step re-probes an exhausted key. Priority keeps the preferred model at the cost of probing; LKGP minimises latency but can stay on a fallback after the preferred model recovers.

C. Evidence-linked notes and independent verification

Does citing evidence reduce unsupported content? With the same Gemini writer and prompt apart from the citation rule, we wrote a single-pass SOAP note and an evidence-linked note for each of the 40 dialogues, shuffled the sentences of both, and had a judge from another vendor (gpt-oss-120b) label each against the *full* transcript without knowing its source (Table II). Evidence-linked notes had fewer unsupported sentences (5.0% against 11.2%, non-overlapping intervals), fewer notes with any unsupported sentence (20/40 against 30/40), and no lower ROUGE or BERTScore.

TABLE II
NOTES ON ACI-BENCH TEST1. TOP: SINGLE-PASS VS EVIDENCE-LINKED NOTES FROM THE SAME WRITER, EVERY SENTENCE JUDGED BLIND AGAINST THE FULL TRANSCRIPT BY ANOTHER VENDOR. BOTTOM: PLANTED ERRORS CAUGHT BY THE VERIFIER, AND FALSE ALARMS ON UNCHANGED SENTENCES.

	Single-pass	Evidence-linked
Sentences per note	20.1	18.2
Supported (95% CI)	80.4% (77.5–83.0)	88.4% (85.9–90.6)
Partly supported	8.3%	6.6%
Unsupported (95% CI)	11.2% (9.2–13.6)	5.0% (3.6–6.8)
Notes with an unsupported sentence	30/40	20/40
ROUGE-1 / 2 / L	44.3 / 16.2 / 41.5	45.4 / 16.9 / 42.5
BERTScore F1	47.1	48.0
	Gemini-written	gpt-oss-written
Changed number caught	95%	95%
Flipped negation caught	100%	100%
Borrowed sentence caught	100%	100%
False alarms on unchanged	7%	12%

Two writers and the cascade. Each of two writers also produced evidence-linked notes judged against only their cited turns, the stricter test the interface uses. Citation coverage was 95.6% (Gemini) and 97.3% (gpt-oss-120b), with 83.9% and 77.0% of sentences supported and 6.9% and 6.7% unsupported. Where the Groq daily cap was reached, verification of Gemini-written notes was cross-model, not cross-vendor (15 notes cross-vendor, 25 cross-model). With $\tau = 0.2$, 1 of 40 Gemini-first notes would have been regenerated.

Does the verifier catch errors? We planted three kinds of error into sentences the verifier had accepted (a changed number, a flipped negation, and a sentence borrowed from another dialogue's note with its citations) and kept other accepted sentences as controls. Across 109 planted errors the verifier flagged 98% as partial or unsupported, with 9% false alarms on 129 unchanged sentences (Table II, bottom); controls had already been accepted, so this measures consistency, not precision on fresh sentences.

Counselling dialogues. On AnnoMI, evidence-linked notes (without sentence repair) for 40 dialogues had 549 sentences, 96.4% cited, of which 80.3% were judged supported, 13.3% partly and 6.4% unsupported. These judgements are not vendor-independent: 39 of 40 notes were checked by a Gemini model other than the writer's, because the Groq quota was exhausted.

D. Role identification

With speaker labels replaced by random identifiers, the eight statistics alone named the clinician correctly in 38 of 40 ACI-Bench dialogues, the model alone in 37, and their combination, as deployed, in 39 (97.5%). Only 2 ACI dialogues had the patient speaking more than the doctor. In AnnoMI counselling the client often talks more: statistics alone were right in 38 of 40 dialogues, the model alone in 40, and the combination in 40; in the 17 dialogues where the client spoke more, the three methods were right in 16, 17 and 17. On PriMock57 the deployed combination was right in 57 of 57 consultations.

TABLE III
SPEECH RESULTS (%). TOP: SIMULATED CONVERSATIONS (12 × 3 CONDITIONS): RECOGNITION WITH AND WITHOUT SEPFORMER DENOISING, AND DER. MIDDLE: SPEAKER-ATTRIBUTED CPWER OF EACH PATH; THE FIRST ROW IS OUR FIRST DESIGN, THE BOLD ROW THE CURRENT ONE. BOTTOM: PRIMOCK57 MOCK CONSULTATIONS.

	Clean	5 dB	0 dB
WER Whisper, original / denoised	2.1 / 2.9	13.5 / 28.9	36.0 / 70.9
WER Gemini, original / denoised	5.9 / 4.0	16.9 / 26.9	45.5 / 65.0
DER pyannote / SpeechBrain	9.3 / 23.1	12.9 / 25.1	22.6 / 24.9
SepFormer + SpeechBrain + Gemini	43.4	48.8	77.9
pyannote + Whisper, by duration	21.4	47.1	68.8
pyannote + Whisper, by word time	**2.6**	**19.7**	**53.2**
SpeechBrain + Whisper, by word time	21.0	27.1	46.9
Cloud audio diarization (Gemini)	17.8	37.1	59.3
PriMock57, 57 consultations, 8.64 h (pyannote + Whisper)			
DER / WER	16.8% / 18.8%		
cpWER by word time / by duration	23.8% / 53.0%		

E. Speech recognition and speaker attribution

Table III reports the speech results. Denoising before recognition, as in our first design, raised word error rates under noise for both recognisers (Whisper at 0 dB from 36.0% to 70.9%; $p < 0.001$ and $p = 0.006$ for Whisper and Gemini), so the deployed pipeline transcribes the original recording. pyannote separated speakers with 9.3% DER on clean audio against 23.1% for SpeechBrain clustering ($p < 0.001$; at 5 dB $p < 0.001$). Speaker attribution depends as much on placing words as on finding turns: spreading the transcript over pyannote's turns by duration gave 21.4% cpWER on clean audio, and assigning each word to the turn containing its Whisper timestamp gave 2.6% ($p < 0.001$ in all three conditions). Against cloud audio diarization the local pipeline attributed better on clean audio ($p = 0.032$) and at 5 dB (19.7% against 37.1%, $p = 0.005$). At 0 dB pyannote's DER rises to 22.6% ($p = 0.345$ against SpeechBrain), and SpeechBrain clustering with word timestamps attributes slightly better (46.9% against 53.2%); the difference from cloud diarization at 0 dB is not significant ($p = 0.520$). On real mock consultations the same pipeline reached 16.8% DER and 18.8% WER, with 23.8% speaker-attributed error against 53.0% when words were spread by duration.

Per minute of audio, pyannote took 2.1 s on a laptop GPU (RTX 4050) and 22.0 s on its CPU, against 78.2 s for the first design's SepFormer enhancement with SpeechBrain diarization; Whisper on Groq took 10.2 s and Gemini speech-to-text 67.3 s.

V. DISCUSSION AND LIMITATIONS

What the results support. Citation-constrained writing halved unsupported content against a single-pass note without lowering overlap with reference notes, and an independent verifier caught most planted factual errors, so sentence-level flags give a reviewer a concrete place to look instead of an unexplained confidence score [47]. Resilience in the router comes from remembering failures. Placing words by timestamp, not turn duration, accounts for most of the gain in speaker attribution on simulated and acted speech, and role identification without enrolment was accurate on medical and counselling dialogue.

Deployment. EchoScribe runs on the clinician's machine; only the configured model vendors receive audio or text, an institution can restrict processing to vendors its governance policy approves, and long recordings are compacted before reaching a cloud model.

Limitations. (i) ACI-Bench and PriMock57 are medical, and AnnoMI has no reference notes; no public corpus of counselling or mentoring sessions with reference notes exists, so GROW notes were not evaluated. (ii) Support labels come from a model; the planted-error test bounds how often it misses errors of three kinds, but subtler omissions and misattributions need human review [7], and where a quota forced a judge from the writer's vendor the check was not vendor-independent. (iii) Simulated speech is read without overlap, and PriMock57 is acted speech; overlapping, emotional clinical speech will be harder, and at 0 dB our attribution is not better than cloud diarization. (iv) Results depend on specific model versions that vendors retire; we log the serving model of every request. (v) The risk screen has not been validated against clinician labels. Blinded rating kits for notes and verifier agreement are released with the code. **Ethics:** no human-subject data were used; all experiments use public datasets and scripted or synthetic material, and free-tier model vendors may use submitted text for training, so deployments with real sessions should use institutionally approved vendors or local models.

VI. CONCLUSION

EchoScribe produces SOAP and GROW notes from recorded two-person sessions in which every sentence points to the turns that support it and is checked by a model other than its writer, with regeneration of failing notes, a two-model risk screen that surfaces disagreement, and an embedded router that keeps every processing role available through quota exhaustion and outages. Citing evidence halved unsupported content against a single-pass note from the same writer; assigning words to speakers by timestamp cut speaker-attributed error from 53.0% to 23.8% on mock consultations; and role identification needed no enrolment. The code, evaluation scripts and rating kits are available at https://github.com/Aravind-Krishnan-S/ECHOSCRIBE.

ACKNOWLEDGMENT

The authors thank the faculty of the Department of Artificial Intelligence and Data Science at Jyothi Engineering College for their guidance, and the pyannote, SpeechBrain, ACI-Bench, AnnoMI, PriMock57, LibriSpeech and OmniRoute communities for openly licensed models, data and code. Claude (Anthropic) was used to draft and revise the manuscript text and LaTeX, and to write parts of the software and evaluation code; the authors directed, verified and take responsibility for all content.

REFERENCES

[1] L. L. Weed, "Medical records that guide and teach," *N. Engl. J. Med.*, vol. 278, no. 11, pp. 593–600, 1968.

[2] J. Whitmore, *Coaching for Performance: The Principles and Practice of Coaching and Leadership*, 5th ed. London, U.K.: Nicholas Brealey, 2017.

[3] C. Sinsky *et al.*, "Allocation of physician time in ambulatory practice: A time and motion study in 4 specialties," *Ann. Intern. Med.*, vol. 165, no. 11, pp. 753–760, 2016.

[4] A. A. Tierney *et al.*, "Ambient artificial intelligence scribes to alleviate the burden of clinical documentation," *NEJM Catal. Innov. Care Deliv.*, vol. 5, no. 3, 2024.

[5] P. J. Lukac, W. Turner, S. Vangala, A. T. Chin, J. Khalili, Y. T. Shih, C. Sarkisian, E. M. Cheng, and J. N. Mafi, "Ambient AI scribes in clinical practice: A randomized trial," *NEJM AI*, vol. 2, no. 12, 2025.

[6] Z. Ji *et al.*, "Survey of hallucination in natural language generation," *ACM Comput. Surv.*, vol. 55, no. 12, 2023.

[7] E. Asgari *et al.*, "A framework to assess clinical safety and hallucination rates of LLMs for medical text summarisation," *npj Digit. Med.*, vol. 8, p. 274, 2025.

[8] H. Wang *et al.*, "An evaluation framework for ambient digital scribing tools in clinical applications," *npj Digit. Med.*, vol. 8, p. 358, 2025.

[9] Y. Xu, H. Jia, M. Wang, J. Feng *et al.*, "Enhancing clinical documentation with voice processing and large language models: A study on the LAOS system," *npj Digit. Med.*, vol. 8, p. 798, 2025.

[10] W.-W. Yim, Y. Fu, A. Ben Abacha, N. Snider, T. Lin, and M. Yetisgen, "ACI-BENCH: A novel ambient clinical intelligence dataset for benchmarking automatic visit note generation," *Sci. Data*, vol. 10, p. 586, 2023.

[11] OmniRoute contributors, "OmniRoute: Open-source AI gateway with combo routing and provider cooldowns, v3.8.52," https://github.com/diegosouzapw/OmniRoute, 2026, mIT License.

[12] A. Ben Abacha, W.-W. Yim, G. Adams, N. Snider, and M. Yetisgen, "Overview of the MEDIQA-Chat 2023 shared tasks on the summarization & generation of doctor-patient conversations," in *Proc. 5th Clinical Natural Language Processing Workshop*, 2023, pp. 503–513.

[13] A. Papadopoulos Korfiatis, F. Moramarco, R. Sarac, and A. Savkov, "PriMock57: A dataset of primary care mock consultations," in *Proc. 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, Dublin, Ireland, 2022, pp. 588–598.

[14] Z. Wu, S. Balloccu, V. Kumar, R. Helaoui, E. Reiter, D. Reforgiato Recupero, and D. Riboni, "Anno-MI: A dataset of expert-annotated counselling dialogues," in *Proc. IEEE ICASSP*, 2022, pp. 6177–6181.

[15] D. Van Veen *et al.*, "Adapted large language models can outperform medical experts in clinical text summarization," *Nat. Med.*, vol. 30, no. 4, pp. 1134–1142, 2024.

[16] H. Wang *et al.*, "Towards adapting open-source large language models for expert-level clinical note generation," in *Findings Assoc. Comput. Linguistics: ACL 2025*, 2025, pp. 12 084–12 117.

[17] P. K. Adhikary *et al.*, "Exploring the efficacy of large language models in summarizing mental health counseling sessions: Benchmark study," *JMIR Ment. Health*, vol. 11, p. e57306, 2024.

[18] T. Gao, H. Yen, J. Yu, and D. Chen, "Enabling large language models to generate text with citations," in *Proc. EMNLP*, 2023, pp. 6465–6488.

[19] S. Min *et al.*, "FActScore: Fine-grained atomic evaluation of factual precision in long form text generation," in *Proc. EMNLP*, 2023, pp. 12 076–12 100.

[20] P. Manakul, A. Liusie, and M. J. F. Gales, "SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models," in *Proc. EMNLP*, 2023, pp. 9004–9017.

[21] S. Dhuliawala *et al.*, "Chain-of-verification reduces hallucination in large language models," in *Findings Assoc. Comput. Linguistics: ACL 2024*, 2024, pp. 3563–3578.

[22] T. J. Park *et al.*, "A review of speaker diarization: Recent advances with deep learning," *Comput. Speech Lang.*, vol. 72, 2022.

[23] B. Desplanques, J. Thienpondt, and K. Demuynck, "ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification," in *Proc. Interspeech*, 2020, pp. 3830–3834.

[24] U. von Luxburg, "A tutorial on spectral clustering," *Stat. Comput.*, vol. 17, no. 4, pp. 395–416, 2007.

[25] A. Plaquet and H. Bredin, "Powerset multi-class cross entropy loss for neural speaker diarization," in *Proc. Interspeech*, 2023, pp. 3222–3226.

[26] T. Park *et al.*, "Sortformer: A novel approach for permutation-resolved speaker supervision in speech-to-text systems," in *Proc. ICML*, ser. PMLR, vol. 267, 2025.

[27] M. Bain, J. Huh, T. Han, and A. Zisserman, "WhisperX: Time-accurate speech transcription of long-form audio," in *Proc. Interspeech*, 2023.

[28] A. Zolensky, K. J. Jang, J. Sabin *et al.*, "Speaker role identification in clinical conversations," medRxiv preprint, 2025.

[29] C. Subakan, M. Ravanelli, S. Cornell, M. Bronzi, and J. Zhong, "Attention is all you need in speech separation," in *Proc. IEEE ICASSP*, 2021, pp. 21–25.

[30] M. Ravanelli *et al.*, "SpeechBrain: A general-purpose speech toolkit," arXiv:2106.04624, 2021.

[31] G. Wichern *et al.*, "WHAM!: Extending speech separation to noisy environments," in *Proc. Interspeech*, 2019.

[32] H. Dubey *et al.*, "ICASSP 2023 deep noise suppression challenge," *IEEE Open J. Signal Process.*, vol. 5, pp. 725–737, 2024.

[33] K. Iwamoto, T. Ochiai, M. Delcroix *et al.*, "How bad are artifacts?: Analyzing the impact of speech enhancement errors on ASR," in *Proc. Interspeech*, 2022.

[34] I. Ong *et al.*, "RouteLLM: Learning to route LLMs with preference data," arXiv:2406.18665, 2024.

[35] L. Chen, M. Zaharia, and J. Zou, "FrugalGPT: How to use large language models while reducing cost and improving performance," *Trans. Mach. Learn. Res.*, 2024.

[36] Gemini Team, Google, "Gemini: A family of highly capable multimodal models," arXiv:2312.11805, 2023.

[37] OpenAI, "gpt-oss-120b & gpt-oss-20b model card," arXiv:2508.10925, 2025.

[38] A. Radford *et al.*, "Robust speech recognition via large-scale weak supervision," in *Proc. ICML*, 2023.

[39] H. Bredin, "pyannote.audio 2.1 speaker diarization pipeline: principle, benchmark, and recipe," in *Proc. Interspeech*, 2023, pp. 1983–1987.

[40] K. Posner *et al.*, "The Columbia-Suicide Severity Rating Scale: Initial validity and internal consistency findings from three multisite studies with adolescents and adults," *Am. J. Psychiatry*, vol. 168, no. 12, pp. 1266–1277, 2011.

[41] K. Kroenke, R. L. Spitzer, and J. B. W. Williams, "The PHQ-9: Validity of a brief depression severity measure," *J. Gen. Intern. Med.*, vol. 16, no. 9, pp. 606–613, 2001.

[42] V. Panayotov, G. Chen, D. Povey, and S. Khudanpur, "Librispeech: An ASR corpus based on public domain audio books," in *Proc. IEEE ICASSP*, 2015, pp. 5206–5210.

[43] C.-Y. Lin, "ROUGE: A package for automatic evaluation of summaries," in *Text Summarization Branches Out (ACL Workshop)*, 2004, pp. 74–81.

[44] T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi, "BERTScore: Evaluating text generation with BERT," in *Proc. ICLR*, 2020.

[45] H. Bredin, "pyannote.metrics: A toolkit for reproducible evaluation, diagnostic, and error analysis of speaker diarization systems," in *Proc. Interspeech*, 2017.

[46] S. Watanabe *et al.*, "CHiME-6 challenge: Tackling multispeaker speech recognition for unsegmented recordings," in *Proc. 6th Int. Workshop on Speech Processing in Everyday Environments (CHiME)*, 2020.

[47] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Proc. ICML*, 2017, pp. 1321–1330.